


Chapter 7: Deadlocks

Narzu Tarannum
BRACU





Covered in chapter 7

- The Deadlock Problem
- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock



Inspiring Excellence





Chapter Objectives

- To develop a description of **deadlocks**, which prevent sets of concurrent processes from completing their tasks.
- To present a number of different methods for **preventing or avoiding deadlocks** in a computer system.




Inspiring Excellence





Bridge Crossing Example

Deadlock cause (a situation or opposing parties) to come to a point where no progress can be made because of fundamental disagreement.



- Each section of a bridge can be viewed as a resource.
- Traffic moves only in one direction on the Bridge.
- Deadlock occurs!
- If deadlock occurs, it can be resolved if one car backs up.
- Several cars may have to be backed up.
- **Deadlock is a situation where a set of processes are blocked because each process is holding a resource and waiting for another resource acquired by some other process.**




Inspiring Excellence





The Deadlock Problem

- In a multiprogramming environment, several processes may compete for a finite number of resources.
- A process requests resources; if the resources are not available at that time, the process enters a wait state. It may happen that waiting process will never gain change state, because the resources they have requested are held by other waiting processes. This situation is called a deadlock.



$P1 \rightarrow R1$ **Need R2**
 $P2 \rightarrow R2$ **Need R1**

- Example

- System has 2 Resources: R1, R2.
- *P1* and *P2* each hold one resource and each needs another one.



Inspiring Excellence





System Model

- A system consist of a **finite** number of resources to be distributed among a number of competing processes.
 - Resource types R_1, R_2, \dots, R_m
 - Process are P_1, P_2, \dots, P_n
- Resources are partitioned into several types
 - *Physical resource for example, CPU cycles, memory space, I/O devices .*
 - Logical resources for example, semaphores, mutex locks, and files.
- Each resource type consisting of some number of **identical** instance.
- Each resource type R_i has W_i instances.



Inspiring Excellence





System Model

- Under the normal mode of operation, a process may utilize a resource in only the following sequence:
 - **Request:** The process requests the resource.
 - **Use:** The process can operate on the resource
 - **Release:** The process releases the resource.
- **A process must request a resource before using it and must release the resource after using it.**



Inspiring Excellence





Necessary Conditions for Deadlock

Deadlock can arise if four conditions hold **simultaneously** in a system:

1. Mutual exclusion
2. Hold and wait
3. No preemption
4. Circular wait

If one of them is not present in a system, no deadlock will arise



Inspiring Excellence








Necessary Conditions

1. Mutual exclusion:

At least one resource must be held in a non-sharable mode; only one process at a time can use a resource.




Inspiring Excellence





Necessary Conditions

2. Hold and wait: a process holding at least one resource is waiting to acquire additional resources held by other processes.



Inspiring Excellence





Necessary Conditions

3. No preemption:

a resource can be released only voluntarily by the process holding it, after that process has completed its task.



Inspiring Excellence






Necessary Conditions

4. Circular wait:

There must be a circular chain of two or more processes each of which is waiting for a resource held by the next member of the chain.



Inspiring Excellence






Resource-Allocation Graph

A set of vertices V and a set of edges E .

- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system.
 $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system.
- request edge – directed edge $P_i \rightarrow R_j$
- assignment edge – directed edge $R_j \rightarrow P_i$




Inspiring Excellence





Resource-Allocation Graph (Cont.)


- Process




- Resource Type with 4 instances



- P_i requests instance of R_j



- P_i is holding an instance of R_j



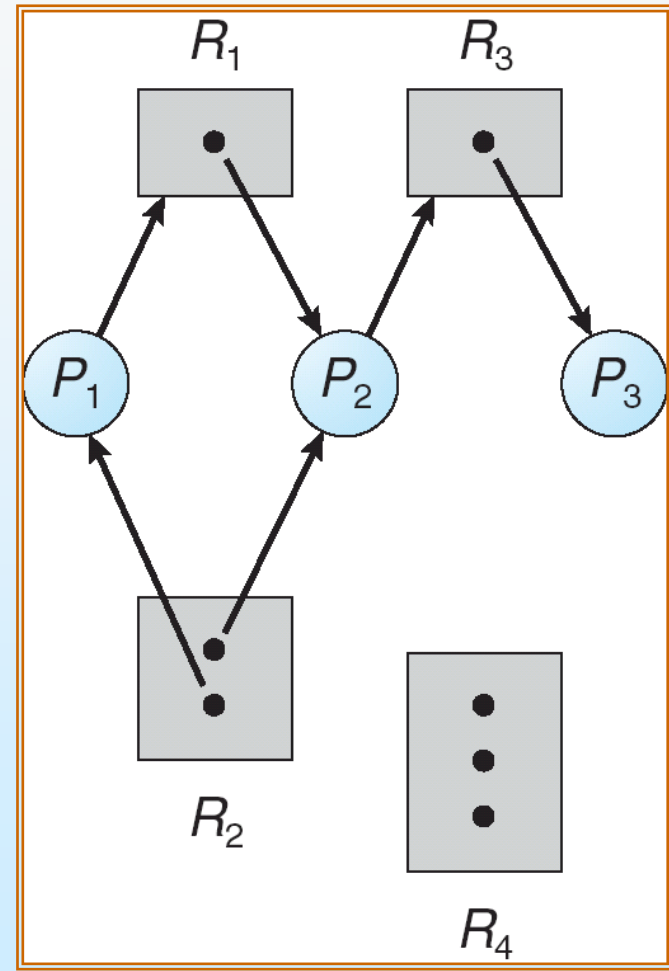
Inspiring Excellence





Example of a Resource Allocation Graph

- $P = \{P_1, P_2, P_3\}$
- $R = \{R_1, R_2, R_3, R_4\}$
- $E = \{P_1 \rightarrow R_1, P_2 \rightarrow R_3, R_1 \rightarrow P_2, R_2 \rightarrow P_2, R_2 \rightarrow P_1, R_3 \rightarrow P_3\}$



Inspiring Excellence





Basic Facts

- If graph contains no cycles \Rightarrow no deadlock.
- If graph contains a cycle \Rightarrow a deadlock may occur
 - if only one instance per resource type, then **deadlock**.
 - if several instances per resource type, **possibility** of deadlock.




Inspiring Excellence





Resource Allocation Graph With A Deadlock



- In the figure At this point, two minimal cycles exist in the system:

- $P1 \rightarrow R1 \rightarrow P2 \rightarrow R3 \rightarrow P3 \rightarrow R2 \rightarrow P1$
- $P2 \rightarrow R3 \rightarrow P3 \rightarrow R2 \rightarrow P2$




Inspiring Excellence





Resource Allocation Graph With A Cycle But No Deadlock



Inspiring Excellence





Methods for Handling Deadlocks

- Ensure that the system will *never* enter a deadlock state (Deadlock prevention).
- In deadlock avoidance, the request for any resource will be granted if the resulting state of the system doesn't cause deadlock in the system. In order to avoid deadlocks, the process must tell OS, the maximum number of resources a process can request to complete its execution. (Deadlock avoidance).
- Allow the system to enter a deadlock state and then recover (Deadlock detection and recovery).
- Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX. (Ignorance)



Inspiring Excellence





Methods for Handling Deadlocks

- Deadlock prevention
- Deadlock avoidance
- Deadlock detection and recovery
- Ignorance



Inspiring Excellence






Deadlock Prevention

Eliminate one of the four conditions.

1. Mutual Exclusion –

- The mutual exclusion condition must hold for non-sharable resources. Example: a printer
- sharable resources, in contrast do not require mutually exclusive access and thus can not be involved in a deadlock. Example: read only files



In general, however, we cannot prevent deadlocks by denying the mutual-exclusion condition, because some resources are intrinsically non-shareable.



Inspiring Excellence





Deadlock Prevention (Cont.)

2. Hold and Wait – must guarantee that *whenever a process requests a resource, it does not hold any other resources.*

- **First Protocol or conservative approach:** A process is allowed to start execution if and only if it has acquired all the resources.
- **Second Protocol or do not hold approach:** A process will acquire only desired resources but before making any fresh request, it must release all the resources that it currently hold.
 - Both this protocols have two main disadvantages:
 - 1. Poor resource utilization
 - 2. starvation is possible
- **Third Protocol or wait-time out approach:** We place a maximum time bound up to which a process can wait for resources, after which they must release all the holding resources.



Inspiring Excellence

R1, R2, ... R4, R5,R10





Deadlock Prevention (Cont.)

3.No Preemption –

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it (that is, the process must wait), then all resources currently being held are released or preempted.
 - Preempted resources are added to the list of resources for which the process is waiting.
 - Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.
- Alternatively, if a process requests some resources, we first check whether they are available. If they are, we allocate them. If they are not, we check whether they are allocated to some other process that is waiting for additional resources. If so, we preempt the desired resources from the waiting process and allocate them to the requesting process.

This protocol is often applied to resources whose state can be easily saved and restored later, such as CPU registers and memory space. It can not generally be applied to such resources as printers and tape drivers.




Inspiring Excellence






Deadlock Prevention (Cont.)

4. Circular Wait – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration.



P1	P2
R1	R2
R2	R1

P1	P2
R1	R1 ✓
R2	R2



Inspiring Excellence






Deadlock Prevention (Cont.)

4. Circular Wait – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration. Let $R = R_1, R_2, \dots, R_m$ are the resources.

-Circular wait can be eliminated by just giving the natural number of every resource. So, we can define the function $F: R \rightarrow N$



Developing and ordering, or hierarchy, in itself does not prevent deadlock. It is up to application developer to write programs that follow the ordering.



Inspiring Excellence





Methods for Handling Deadlocks

- Deadlock prevention
- Deadlock avoidance
- Deadlock detection and recovery
- Ignorance



Inspiring Excellence






Deadlock Avoidance

An alternative method for avoiding deadlocks is to require additional information about how resources are to be requested.

- Banker's Algorithm requires the information of each process that declare the *maximum number* of resources of each type that it may need.
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition.
- Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes.

With this knowledge of the complete sequence of requests and releases for each process, the system can decide for each request whether or not the process should wait in order to avoid a possible future deadlock.



Inspiring Excellence





Safe-Unsafe State

- A state is safe if the system can allocate resources to each process (up to its maximum) in some order and still avoid a deadlock.
 - A state is said to be safe, if there is some scheduling order in which every process can run to completion that means system is in safe state **if there exists a safe sequence of all processes.**
 - An unsafe state **does not have to** lead to a deadlock; it **could** lead to a deadlock.

ENSURE SYSTEM NEVER REACHES AN UNSAFE STATE




Inspiring Excellence





Safe, Unsafe , Deadlock State

- If a system is in safe state \Rightarrow no deadlocks.
- If a system is in unsafe state \Rightarrow possibility of deadlock. Not all unsafe states are deadlocks, however.
- Avoidance \Rightarrow ensure that a system will never enter an unsafe state.



Inspiring Excellence





Example of a Safe-unsafe state

Let there are 3 processes (A, B, C) where total number of resources are 10. Every process has to declare maximum number of resource they need in advance. In present scenario system allocated 3 units of resources to A where maximum requirement of A is 9 units. In the same way system allocated 2 units to B where maximum requirement of B is 4 and system allocated 2 units to C where maximum requirement of C is 7

Total: 10 units

	Has	Max
A	3	9
B	2	4
C	2	7

Free: 3 units

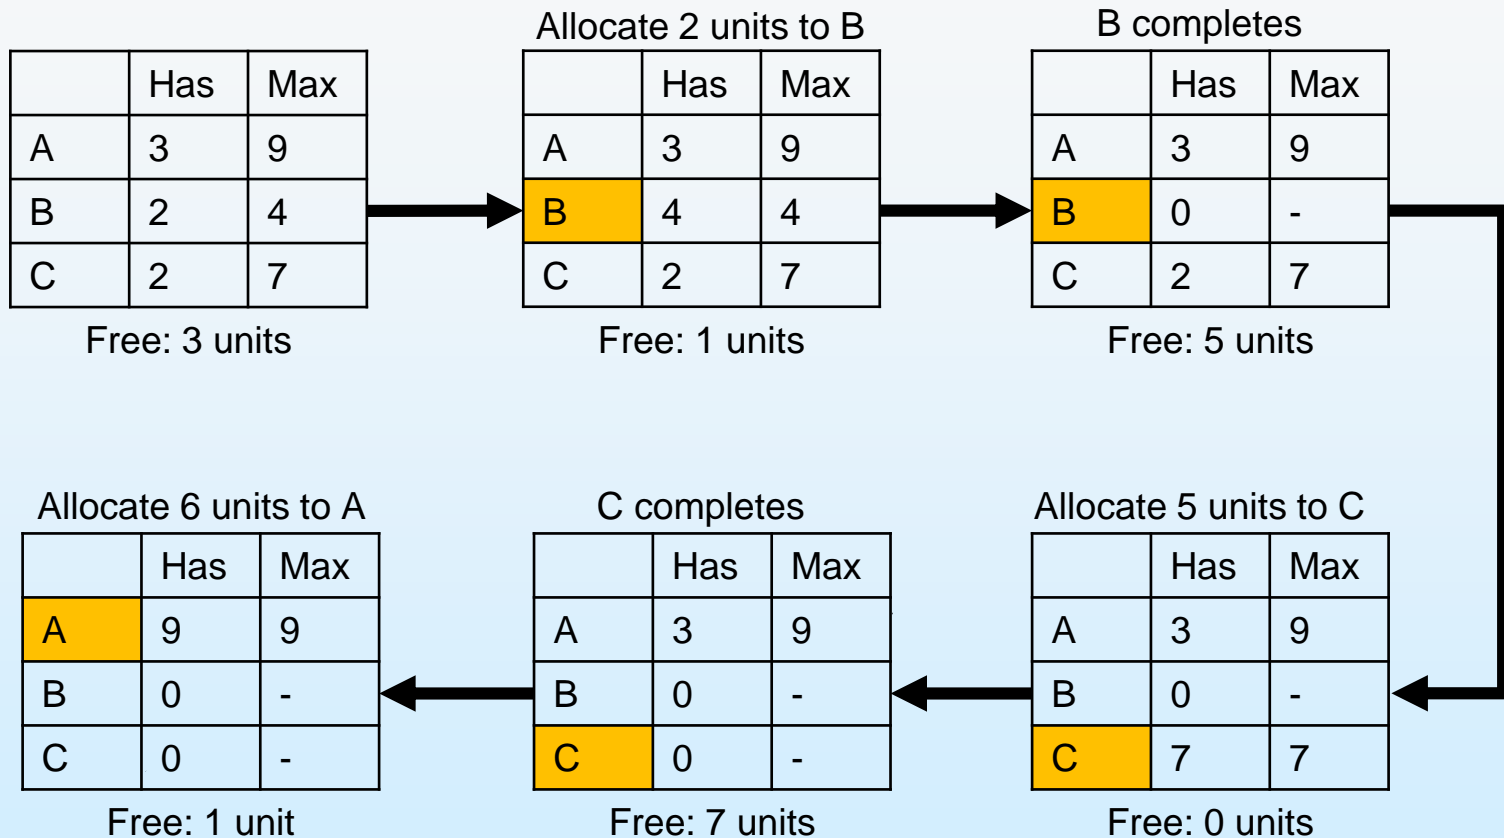


Inspiring Excellence





Safe state

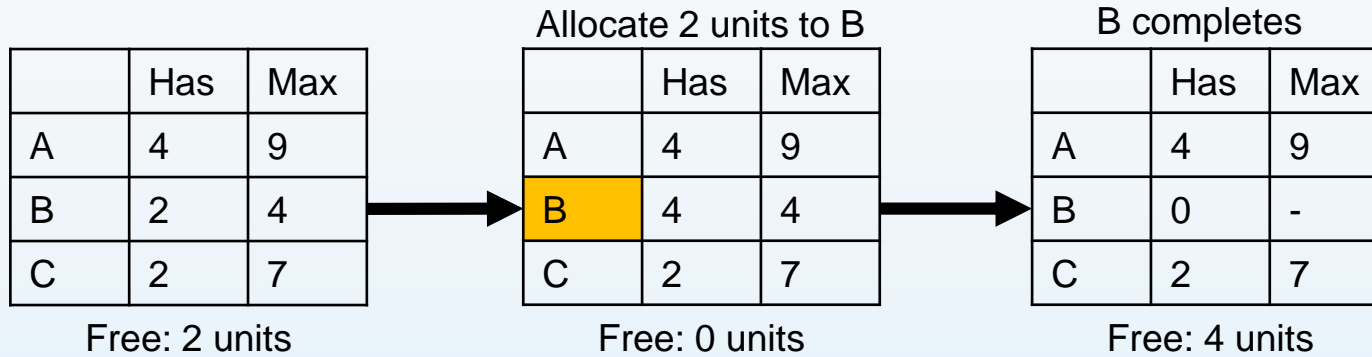


This is a safe state because there is some scheduling order in which every process executes. Here the order is: $B \rightarrow C \rightarrow A$





Unsafe state



5 units to C or A **can not be allocated**

	Has	Max
A	?	9
B	0	-
C	?	7

Free: 0 units



Inspiring Excellence

This is an unsafe state because there exist no scheduling order in which every process executes.





Banker's Algorithm

- Multiple instances of each resource type.
- Each process must a priori claim maximum use.
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.



Inspiring Excellence





Data Structures for the Banker's Algorithm

Let n = number of processes, and m = number of resources types.

- **Available:** Vector of length m . If available $[j] = k$, there are k instances of resource type R_j available.
- **Max:** $n \times m$ matrix. If $Max[i, j] = k$, then process P_i may request at most k instances of resource type R_j .
- **Allocation:** $n \times m$ matrix. If $Allocation[i, j] = k$ then P_i is currently allocated k instances of R_j .
- **Need:** $n \times m$ matrix. If $Need[i, j] = k$, then P_i may need k more instances of R_j to complete its task.

$$Need[i, j] = Max[i, j] - Allocation[i, j].$$

$$X \leq Y \quad \text{if } X=(1,7,3,2) \text{ and } Y=(0,3,2,1) \text{ then } X \neq Y.$$



Inspiring Excellence





Banker's Safety Algorithm

1. Let Work and Finish be vectors of length m and n, respectively.
Initialize:
 $Work = Available$
 $Finish[i] = false$ for $i=0, 1, \dots, n-1$.
2. Find an i such that both:
(a) $Finish[i] = false$
(b) $Need_i \leq Work$
If no such i exists, go to step 4.
3. $Work = Work + Allocation_i$
 $Finish[i] = true$
go to step 2.
4. If $Finish[i] == true$ for all i, then the system is in a safe state.



Inspiring Excellence





Example of Banker's Algorithm

- 3 resource types:
 A (10 instances), B (5 instances), and C (7 instances).
- 5 processes P_0 through P_4 ;
- Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>	<u>Need = Max - Allocation</u>
	$A\ B\ C$	$A\ B\ C$	$A\ B\ C$	$A\ B\ C$
P_0	0 1 0	7 5 3	3 3 2	7 4 3
P_1	2 0 0	3 2 2		1 2 2
P_2	3 0 2	9 0 2		6 0 0
P_3	2 1 1	2 2 2		0 1 1
P_4	0 0 2	4 3 3		4 3 1



Inspiring Excellence





Example of Banker's safety Algorithm

- Let Work and Finish be vectors of length m and n, respectively.
Initialize:
Work = Available
Finish[i] = false for $i=0,1,\dots, n-1$
- Find an i such that both:
(a) Finish[i] = false
(b) $Need_i \leq Work$
If no such i exists, go to step 4.
- Work = Work + Allocation_i
Finish[i] = true
go to step 2.
- If Finish[i] == true for all i, then the system is in a safe state.

Process	Allocation	Max	Available	Need
	A B C	A B C	A B C	A B C
P_0	0 1 0	7 5 3	3 3 2	7 4 3
P_1	2 0 0	3 2 2		1 2 2
P_2	3 0 2	9 0 2		6 0 0
P_3	2 1 1	2 2 2		0 1 1
P_4	0 0 2	4 3 3		4 3 1

Work = Available = 3 3 2

Finish[0]=Finish[1]=Finish[2]=Finish[3]=Finish[4]=False

Need(P_0) > Work \rightarrow Finish[0]=F \rightarrow Do Nothing	$7\ 4\ 3 \leq 3\ 3\ 2 \rightarrow F \rightarrow$ do nothing
Need(P_1) \leq Work \rightarrow Finish[1]=T \rightarrow Work=Work + Allocation P_1	$1\ 2\ 2 \leq 3\ 3\ 2 \rightarrow T$: Work = $3\ 3\ 2 + 2\ 0\ 0 = 5\ 3\ 2$
Need(P_2) > Work \rightarrow Finish[2]=F \rightarrow Do Nothing	$6\ 0\ 0 \leq 5\ 3\ 2 \rightarrow F \rightarrow$ do nothing
Need(P_3) \leq Work \rightarrow Finish[3]=T \rightarrow Work=Work + Allocation P_3	$0\ 1\ 1 \leq 5\ 3\ 2 \rightarrow T$: Work = $5\ 3\ 2 + 2\ 1\ 1 = 7\ 4\ 3$
Need(P_4) \leq Work \rightarrow Finish[4]=T \rightarrow Work=Work + Allocation P_4	$4\ 3\ 1 \leq 7\ 4\ 3 \rightarrow T$: Work = $7\ 4\ 3 + 0\ 0\ 2 = 7\ 4\ 5$
Need(P_0) \leq Work \rightarrow Finish[0]=T \rightarrow Work=Work + Allocation P_0	$7\ 4\ 3 \leq 7\ 4\ 5 \rightarrow T$: Work = $7\ 4\ 5 + 0\ 1\ 0 = 7\ 5\ 5$
Need(P_2) \leq Work \rightarrow Finish[2]=T \rightarrow Work=Work + Allocation P_2	$6\ 0\ 0 \leq 7\ 5\ 5 \rightarrow T$: Work = $7\ 5\ 5 + 3\ 0\ 2 = 10\ 5\ 7$



Inspiring Excellence

Found the safe sequence <P1, P3, P4, P0, P2>





- Request for resources (1 0 2) for P1 will be granted at this state?



Inspiring Excellence





Banker's Resource-request Algorithm

Request = request vector for process P_i . If $\text{Request}_i[j] = k$ then process P_i wants k instances of resource type R_j .

$P_i \rightarrow \text{Request}_i$

1. If $\text{Request}_i \leq \text{Need}_i$ go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If $\text{Request}_i \leq \text{Available}$, go to step 3. Otherwise P_i must wait, since resources are not available.
3. Pretend to allocate requested resources to P_i by modifying the state as follows:

$\text{Available} = \text{Available} - \text{Request}_i;$

$\text{Allocation}_i = \text{Allocation}_i + \text{Request}_i;$

$\text{Need}_i = \text{Need}_i - \text{Request}_i;$

4. Check Bankers safety Algorithm, if this new state is safe

If safe \Rightarrow the resources are allocated to P_i .

If unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored



Inspiring Excellence





Example of Banker's Resource-request Algorithm

Request = request vector for process P_i . If $\text{Request}_i[j] = k$ then process P_i wants k instances of resource type R_j .

1. If $\text{Request}_i \leq \text{Need}_i$ go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.

2. If $\text{Request}_i \leq \text{Available}$, go to step 3. Otherwise P_i must wait, since resources are not available.

3. Pretend to allocate requested resources to P_i by modifying the state as follows:

$\text{Available} = \text{Available} - \text{Request}_i$;

$\text{Allocation}_i = \text{Allocation}_i + \text{Request}_i$;

$\text{Need}_i = \text{Need}_i - \text{Request}_i$;

Process	Allocation	Max	Available	Need
	A B C	A B C	A B C	A B C
P_0	0 1 0	7 5 3	3 3 2	7 4 3
P_1	2 0 0	3 2 2		1 2 2
P_2	3 0 2	9 0 2		6 0 0
P_3	2 1 1	2 2 2		0 1 1
P_4	0 0 2	4 3 3		4 3 1

Check a new additional request for P_1 (1 0 2) can be granted now?

1. if Request of $P_1 \leq \text{Need of } P_1$	$1\ 0\ 2 \leq 1\ 2\ 2 \rightarrow T$
2. if Request of $P_1 \leq \text{Available}$	$1\ 0\ 2 \leq 3\ 3\ 2 \rightarrow T$
3. Available = Available – Request; Allocation _i = Allocation _i + Request; Need _i = Need _i – Request;	Available = $3\ 3\ 2 - 1\ 0\ 2 = 2\ 3\ 0$ Allocation = $2\ 0\ 0 + 1\ 0\ 2 = 3\ 0\ 2$ Need = $1\ 2\ 2 - 1\ 0\ 2 = 0\ 2\ 0$



Inspiring Excellence





Example of Banker's Resource-request Algorithm

3. Available = $3\ 3\ 2 - 1\ 0\ 2 = 2\ 3\ 0$
Allocation = $2\ 0\ 0 + 1\ 0\ 2 = 3\ 0\ 2$
Need = $1\ 2\ 2 - 1\ 0\ 2 = 0\ 2\ 0$

Update table after additional request for P_1 (1 0 2)

Process	Allocation	Max	Available	Need
	A B C	A B C	A B C	A B C
P_0	0 1 0	7 5 3	3 3 2	7 4 3
P_1	2 0 0	3 2 2		1 2 2
P_2	3 0 2	9 0 2		6 0 0
P_3	2 1 1	2 2 2		0 1 1
P_4	0 0 2	4 3 3		4 3 1

Process	Allocation	Max	Available	Need
	A B C	A B C	A B C	A B C
P_0	0 1 0	7 5 3	2 3 0	7 4 3
P_1	3 0 2	3 2 2		0 2 0
P_2	3 0 2	9 0 2		6 0 0
P_3	2 1 1	2 2 2		0 1 1
P_4	0 0 2	4 3 3		4 3 1



Inspiring Excellence





Verify with Banker's safety Algorithm

- Let Work and Finish be vectors of length m and n, respectively.
Initialize:
Work = Available
Finish[i] = false for $i=0,1,\dots, n-1$
- Find an i such that both:
(a) Finish[i] = false
(b) $Need_i \leq Work$
If no such i exists, go to step 4.
- Work = Work + Allocation_i
Finish[i] = true
go to step 2.
- If Finish[i] == true for all i, then the system is in a safe state.

Process	Allocation	Max	Available	Need
	A B C	A B C	A B C	A B C
P_0	0 1 0	7 5 3	2 3 0	7 4 3
P_1	3 0 2	3 2 2		0 2 0
P_2	3 0 2	9 0 2		6 0 0
P_3	2 1 1	2 2 2		0 1 1
P_4	0 0 2	4 3 3		4 3 1

Need(P0) > Work → Finish[1]=F → Do Nothing	7 4 3 ≤ 2 3 0 → F → do nothing
Need(P1) ≤ Work → Finish[1]=T → Work=Work + Allocation	0 2 0 ≤ 2 3 0 → T: Work = 2 3 0 + 3 0 2 = 5 3 2
Need(P2) > Work → Finish[2]=F → Do Nothing	6 0 0 ≤ 5 3 2 → F → do nothing
Need(P3) ≤ Work → Finish[3]=T → Work=Work + Allocation	0 1 1 ≤ 5 3 2 → T: Work = 5 3 2 + 2 1 1 = 7 4 3
Need(P4) ≤ Work → Finish[4]=T → Work=Work + Allocation	4 3 1 ≤ 7 4 3 → T: Work = 7 4 3 + 0 0 2 = 7 4 5
Need(P0) ≤ Work → Finish[0]=T → Work=Work + Allocation	7 4 3 ≤ 7 4 5 → T: Work = 7 4 5 + 0 1 0 = 7 5 5
Need(P2) ≤ Work → Finish[2]=T → Work=Work + Allocation	6 0 0 ≤ 7 5 5 → T: Work = 7 5 5 + 3 0 2 = 10 5 7



Inspiring Excellence

Since there is a safe sequence **<P1, P3, P4, P0, P2>** ,
Additional request P_1 (1 0 2) can be granted





Are following requests get accepted?

- Request for (3 3 0) resources for P4
- Request for (0 2 0) resources for P0

Process	Allocation	Max	Available	Need
	A B C	A B C	A B C	A B C
P_0	0 1 0	7 5 3	3 3 2	7 4 3
P_1	2 0 0	3 2 2		1 2 2
P_2	3 0 2	9 0 2		6 0 0
P_3	2 1 1	2 2 2		0 1 1
P_4	0 0 2	4 3 3		4 3 1



Inspiring Excellence





Are following requests get accepted?

- Request for (3 3 0) resources for P4
- Request for (0 2 0) resources for P0

Process	Allocation	Max	Available	Need
	A B C	A B C	A B C	A B C
P_0	0 1 0	7 5 3	3 3 2	7 4 3
P_1	2 0 0	3 2 2		1 2 2
P_2	3 0 2	9 0 2		6 0 0
P_3	2 1 1	2 2 2		0 1 1
P_4	0 0 2	4 3 3		4 3 1



Inspiring Excellence





Thank You



Inspiring Excellence

